

# Introduction to R

Presented by:



BROWN  
School of Public Health



# Intro to R Programming for Biostatistics

## Day 2 - Arranging and Summarizing Data in R

Adam J Sullivan

# Arranging the Data

# Arranging the Data

- We also have need to make sure the data is ordered in a certain manner. This can be easily done in R with the `arrange()` function.
- Again we can do this in base R but this is not always a clear path.

# Arranging the Data Example

- Let's say that we wish to look at only carriers and departure delay and we wish to order departure delays from what smallest to largest.
- In base R we would have to run the following command:

```
flights[order(flights$dep_delay), c("carrier", "dep_delay")]
```

# Enter the `arrange()` Function

We could do this in an easy manner using the `arrange()` function:

```
arrange(.data, ...)
```

Where

- `.data` is a data frame of interest.
- `...` are the variables you wish to sort by.

# Arranging the Data Example Continued

```
flights %>%  
  select(carrier, dep_delay) %>%  
  arrange(dep_delay)
```

# Arranging the Data Example Continued

```
## # A tibble: 336,776 × 2
##   carrier dep_delay
##   <chr>      <dbl>
## 1      B6        -43
## 2      DL        -33
## 3      EV        -32
## 4      DL        -30
## 5      F9        -27
## 6      MQ        -26
## 7      EV        -25
## 8      MQ        -25
## 9      9E        -24
## 10     B6        -24
## # ... with 336,766 more rows
```



# Arranging the Data Example Continued

- With `arrange()` we first use `select()` to pick the only columns that we want and then we arrange by the `dep_delay`.
- If we had wished to order them in a descending manner we could have simply used the `desc()` function:

```
flights %>%  
  select(carrier, dep_delay) %>%  
  arrange(desc(dep_delay))
```

# More Complex Arrange

- Lets consider that we wish to look at the top 3 departure delays for each day.
- Then we wish to order them from largest to smallest departure delay.
- We then need to do the following:
  1. Group by month and Day
  2. Pick the top 3 departure delays
  3. order them largest to smallest

--- .class #id

# More Complex Arrange Continued

```
flights %>%  
  group_by(month, day) %>%  
  top_n(3, dep_delay) %>%  
  arrange(desc(dep_delay))
```

Where

- `group_by()` is a way to group data. This way we perform operations on a group. So top 3 delays are by a group of day and month.
- `top_n()` takes a tibble and returns a specific number of rows based on a chosen value.

# More Complex Arrange Continued

```
## Source: local data frame [1,108 x 19]
## Groups: month, day [365]
##
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     9     641           900       1301    1242
## 2  2013     6    15    1432          1935       1137    1607
## 3  2013     1    10    1121          1635       1126    1239
## 4  2013     9    20    1139          1845       1014    1457
## 5  2013     7    22     845          1600       1005    1044
## 6  2013     4    10    1100          1900        960    1342
## 7  2013     3    17    2321           810        911     135
## 8  2013     6    27     959          1900        899    1236
## 9  2013     7    22    2257           759        898     121
## 10 2013    12     5     756          1700        896    1058
## # ... with 1,098 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

## On Your Own: RStudio Practice

- Perform the following operations:
  - Group by month and day.
  - use `sample_n()` to pick 1 observation per day.
  - Arrange by longest to smallest departure delay.

--- .class #id

## On Your Own: RStudio Practice

Your answer ***may*** look like:

```
flights %>%  
  group_by(month, day) %>%  
  sample_n(1) %>%  
  arrange(desc(dep_delay))
```

# Summarizing Data

# Summarizing Data

- As you have seen in your own work, being able to summarize information is crucial.
- We need to be able to take out data and summarize it as well.
- We will consider doing this using the `summarise()` function.

# Summarizing Data

- Like in the rest of these lessons, let's consider what happens when we try to do this in base R. We will:
  1. Create a table grouped by dest.
  2. Summarize each group by taking mean of arr\_delay.

```
head(with(flights, tapply(arr_delay, dest, mean, na.rm=TRUE)))  
head(aggregate(arr_delay ~ dest, flights, mean))
```

# Enter summarise() Function

- The summarise() function is:

```
summarise(.data, ...)
```

- where
  - .data is the tibble of interest.
  - ... is a list of name paired summary functions
  - Such as:
    - mean()
    - median
    - var()
    - sd()
    - min()
    - `max()
    -



# Summarizing Data Example

```
flights %>%  
  group_by(dest) %>%  
  summarise(avg_delay = mean(arr_delay, na.rm=TRUE))
```

# Summarizing Data Example

- Consider the logic here:
  1. Group flights by destination
  2. Find the average delay of the groups and call it `avg_delay`.
- This is much easier to understand than the Base R code.

```
## # A tibble: 105 × 2
##   dest avg_delay
##   <chr>   <dbl>
## 1   ABQ  4.381890
## 2   ACK  4.852273
## 3   ALB 14.397129
## 4   ANC -2.500000
## 5   ATL 11.300113
## 6   AUS  6.019909
## 7   AVL  8.003831
## 8   BDL  7.048544
## 9   BGR  8.027933
## 10  BHM 16.877323
## # ... with 95 more rows
```

## Another Example

- Lets say that we would like to have more than just the averages but we wish to have the minimum and the maximum departure delays by carrier:

```
flights %>%  
  group_by(carrier) %>%  
  summarise_each(funs(min(., na.rm=TRUE), max(., na.rm=TRUE)), matches("delay"))
```

# Another Example

```
## # A tibble: 16 × 5
##   carrier dep_delay_min arr_delay_min dep_delay_max arr_delay_max
##   <chr>         <dbl>         <dbl>         <dbl>         <dbl>
## 1      9E          -24          -68           747           744
## 2      AA          -24          -75          1014          1007
## 3      AS          -21          -74           225           198
## 4      B6          -43          -71           502           497
## 5      DL          -33          -71           960           931
## 6      EV          -32          -62           548           577
## 7      F9          -27          -47           853           834
## 8      FL          -22          -44           602           572
## 9      HA          -16          -70          1301          1272
## 10     MQ          -26          -53          1137          1127
## 11     OO          -14          -26           154           157
## 12     UA          -20          -75           483           455
## 13     US          -19          -70           500           492
## 14     VX          -20          -86           653           676
## 15     WN          -13          -58           471           453
## 16     YV          -16          -46           387           381
```

# On Your Own: RStudio Practice

- The following is a new function:
  - Helper function `n()` counts the number of rows in a group
- Then for each day:
  - count total flights
  - Sort in descending order.

# On Your Own: RStudio Practice

Your answer should look like:

```
## Source: local data frame [365 x 3]
## Groups: month [12]
##
##   month   day flight_count
##   <int> <int>      <int>
## 1     11    27         1014
## 2      7    11         1006
## 3      7     8         1004
## 4      7    10         1004
## 5     12     2         1004
## 6      7    18         1003
## 7      7    25         1003
## 8      7    12         1002
## 9      7     9         1001
## 10     7    17         1001
## # ... with 355 more rows
```

We could also have used what is called the `tally()` function:

```
flights %>%
  group_by(month, day) %>%
  tally(sort = TRUE)
```

# Adding New Variables

# Adding New Variables

- There is usually no way around needing a new variable in your data.
- For example, most medical studies have height and weight in them, however many times what a researcher is interested in using is Body Mass Index (BMI).
- We would need to add BMI in.



# Adding New Variables

- Using the tidyverse we can add new variables in multiple ways
  - `mutate()`
  - `transmute()`

# Adding New Variables

With `mutate()` we have

```
mutate(.data, ...)
```

where

- `.data` is your tibble of interest.
- `...` is the name paired with an expression

# Adding New Variables

Then with `transmute()` we have:

```
transmute(.data, ...)
```

where

- `.data` is your tibble of interest.
- `...` is the name paired with an expression

# Differences Between `mutate()` and `transmute()`

- There is only one major difference between `mutate()` and `transmute` and that is what it keeps in your data.
  - `mutate()`
    - creates a new variable
    - It keeps all existing variables
  - `transmute()`
    - creates a new variable.
    - It only keeps the new variables

# Example

- Let's say we wish to have a variable called speed. We want to basically do:

$$\text{speed} = \frac{\text{distance}}{\text{time}} * 60$$

We can first do this with `mutate()`:

```
flights %>%  
  select(flight, distance, air_time) %>%  
  mutate(speed = distance/air_time*60)
```

# Example

```
## # A tibble: 336,776 × 4
##   flight distance air_time    speed
##   <int>    <dbl>    <dbl>    <dbl>
## 1    1545     1400      227 370.0441
## 2    1714     1416      227 374.2731
## 3    1141     1089      160 408.3750
## 4     725     1576      183 516.7213
## 5     461      762      116 394.1379
## 6    1696      719      150 287.6000
## 7     507    1065      158 404.4304
## 8    5708      229       53 259.2453
## 9       79     944      140 404.5714
## 10    301     733      138 318.6957
## # ... with 336,766 more rows
```

# Example

```
flights %>%  
  select(flight, distance, air_time) %>%  
  transmute(speed = distance/air_time*60)
```

# Example

```
flights %>%  
  select(flight, distance, air_time) %>%  
  transmute(speed = distance/air_time*60)
```



#Further Summaries

# Further Summaries

- We have so far discussed how one could find the basic number summaries:
  - mean
  - median
  - standard deviation
  - variance
  - minimum
  - maximum
- However there are many more operations that you may wish to do for summarizing data.
- In fact many of the following examples are excellent choices for working with categorical data which does not always make sense to do the above summaries for.

# Further Summaries

- We will consider:
  1. Grouping and Counting
  2. Grouping, Counting and Sorting
  3. Other Groupings
  4. Counting Groups

# Grouping and Counting

- We have seen the functions `tally()` and `count()`.
- Both of these can be used for grouping and counting.
- They also are very concise in how they are called.

# Grouping and Counting

- For example if we wished to know how many flights there were by month, we would use `tally()` in this manner:

```
flights %>%  
  group_by(month) %>%  
  tally()
```

# Grouping and Counting

- Where as we could do the same thing with `count()`

```
flights %>%  
  count(month)
```

\*Notice: `count()` allowed for month to be called inside of it, removing the need for the `group_by()` function.

# Grouping, counting and sorting.

- Both `tally()` and `count()` have an argument called `sort()`.
- This allows you to go one step further and group by, count and sort at the same time.
- For `tally()` this would be:

```
flights %>% group_by(month) %>% tally(sort=TRUE)
```

```
## # A tibble: 12 × 2
##   month      n
##   <int> <int>
## 1       7 29425
## 2       8 29327
## 3      10 28889
## 4       3 28834
## 5       5 28796
## 6       4 28330
## 7       6 28243
## 8      12 28135
## 9       9 27574
## 10      11 27268
## 11       1 27004
## 12       2 24951
```



- Then for `count()` we would have:

```
flights %>% count_(month, sort=TRUE)
```

- Then for `count()` we would have:

```
## Error in as.lazy_dots(.dots): object 'month' not found
```

# Grouping with other functions

- We can also sum over other values rather than just counting the rows like the above examples.
- For example let us say we were interested in knowing the total distance for planes in a given month.
- We could do this with the `summarise()` function, `tally()` function or the `count()` function:

```
flights %>%  
  group_by(month) %>%  
  summarise(dist = sum(distance))
```

# Grouping with other functions

- We take flights then group by month and then create a new variable called distance, where we sum the distance.
- For `tally()` we could do:

```
flights %>%  
  group_by(month) %>%  
  tally(wt = distance)
```

*Note: in `tally()` the `wt` stands for weight and allows you to weight the sum based on the distance.*

# Grouping with other functions

- With the `count()` function we also use `wt`:

```
flights %>% count(month, wt = distance)
```

```
## # A tibble: 12 × 2
##   month      n
##   <int>   <dbl>
## 1      1 27188805
## 2      2 24975509
## 3      3 29179636
## 4      4 29427294
## 5      5 29974128
## 6      6 29856388
## 7      7 31149199
## 8      8 31149334
## 9      9 28711426
## 10     10 30012086
## 11     11 28639718
## 12     12 29954084
```

## ##Counting Groups

- We may want to know how large our groups are. To do this we can use the following functions:
  - `group_size()` is a function that returns counts of group.
  - `n_groups()` returns the number of groups

## ##Counting Groups

- So if wanted to count the number of flights by month, we could group by month and find the groups size using `group_size()`:

```
flights %>%  
  group_by(month) %>%  
  group_size()
```

##Counting Groups

## [1] 27004 24951 28834 28330 28796 28243 29425 29327 27574 28889 27268  
## [12] 28135



## ##Counting Groups

- If we just wished to know how many months were represented in our data we could use the `n_groups()` function:

```
flights %>%  
  group_by(month) %>%  
  n_groups()
```

##Counting Groups

## [1] 12