

Introduction to R

Presented by:



BROWN
School of Public Health



Intro to R Programming for Biostatistics

Day 3 - Regression in R

Adam J Sullivan
by-nc-nd

Regression Models in R

Linear Regression in R

Linear Regression in R

- We can use R to easily fit linear regressions for us.
- This section will explore the basic commands for linear regression as well as how to test assumptions.
- We will not teach linear regression, but only seek to display how R does it.

lm() Function in R

- To fit Linear Regression models in R we use the `lm()` function.

```
lm(formula, data, subset, weights, na.action,  
    method = "qr", model = TRUE, x = FALSE, y = FALSE, qr = TRUE,  
    singular.ok = TRUE, contrasts = NULL, offset, ...)
```

- `formula` is the regressino equation written as $y \sim x_1 + x_2 + \dots$
- `data` is the dataframe of interest.
- `subset` specific subset of data.
- `weights` for weighted data.
-

Gapminder Data

- Worldwide data source.
- Contains 6 variables

Variable	Meaning
country	Country Name
continent	Continent Name
year	Year Data Accounts For
lifeExp	Life Expectancy at Birth
pop	Total Population
gdpPercap	per-Capita GDP

- Per-capita GDP (Gross domestic product) is given in units of international dollars, "a hypothetical unit of currency that has the same purchasing power parity that the U.S. dollar had in the United States at a given point in time" -- 2005, in this case.

Gapminder Data

```
library(gapminder)
```


Gapminder Data

```
## # A tibble: 1,704 × 6
##   country continent year lifeExp      pop gdpPercap
##   <fctr>    <fctr> <int>   <dbl>    <int>    <dbl>
## 1 Afghanistan      Asia  1952  28.801  8425333  779.4453
## 2 Afghanistan      Asia  1957  30.332  9240934  820.8530
## 3 Afghanistan      Asia  1962  31.997 10267083  853.1007
## 4 Afghanistan      Asia  1967  34.020 11537966  836.1971
## 5 Afghanistan      Asia  1972  36.088 13079460  739.9811
## 6 Afghanistan      Asia  1977  38.438 14880372  786.1134
## 7 Afghanistan      Asia  1982  39.854 12881816  978.0114
## 8 Afghanistan      Asia  1987  40.822 13867957  852.3959
## 9 Afghanistan      Asia  1992  41.674 16317921  649.3414
## 10 Afghanistan     Asia  1997  41.763 22227415  635.3414
## # ... with 1,694 more rows
```

Gapminder Regression

```
kenya <- gapminder %>% filter(country=="Kenya")
kenya_model <- lm(lifeExp ~ year, data=kenya)
kenya_model
```

```
##
## Call:
## lm(formula = lifeExp ~ year, data = kenya)
##
## Coefficients:
## (Intercept)      year
##   -356.1010      0.2065
```

Gapminder Regression

- Basic Model Statement does not include much information
- If we look further we can see what type of object is returned:

```
str(kenya_model)
```

Gapminder Regression

- We see that this is a list
- We can find out what is contained in a list by using the `names()` function.

```
names(kenya_model)
```

```
## [1] "coefficients" "residuals"    "effects"      "rank"
## [5] "fitted.values" "assign"       "qr"          "df.residual"
## [9] "xlevels"      "call"        "terms"       "model"
```

Gapminder Regression

- We see different values that are listed here.
- Lets look at the coefficients

```
kenya_model$coefficients
```

```
## (Intercept)      year
## -356.1009769    0.2065077
```

Gapminder Regression

- We can use other commands on a regression
- For example we can use the `summary()` function:

```
summary(kenya_model)
```

Gapminder Regression

```
##  
## Call:  
## lm(formula = lifeExp ~ year, data = kenya)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -6.3354 -3.5739 -0.2819  3.9984  5.5687   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept) -356.10098   145.08537   -2.454   0.0340 *      
## year          0.20651     0.07329    2.818   0.0182 *      
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 4.382 on 10 degrees of freedom  
## Multiple R-squared:  0.4426, Adjusted R-squared:  0.3868   
## F-statistic: 7.939 on 1 and 10 DF,  p-value: 0.01823
```

Gapminder Regression

```
kenya_summary <- summary(kenya_model)
names(kenya_summary)
```

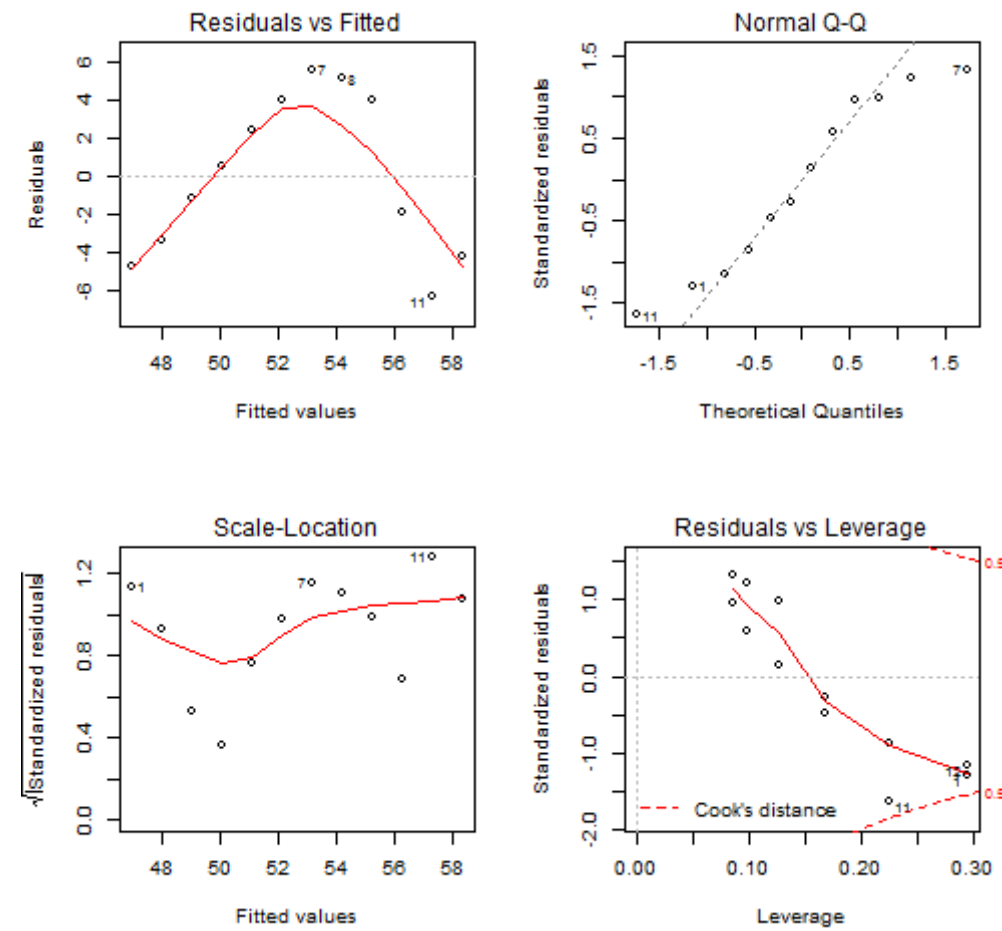

Other Regression Functions

- Other useful functions are listed below:
 - `coefficients(kenya_model)` # model coefficients
 - `confint(kenya_model, level=0.95)` # CIs for model parameters
 - `fitted(kenya_model)` # predicted values
 - `residuals(kenya_model)` # residuals
 - `anova(kenya_model)` # anova table
 - `vcov(kenya_model)` # covariance matrix for model parameters
 - `influence(kenya_model)` # regression diagnostics

Diagnostic Model Plots

```
par(mfrow=c(2,2)) # optional 4 graphs/page  
plot(kenya_model)
```

Diagnostic Model Plots



Comparing Models

- We can compare nested models using the `anova()` function.

```
kenya_model2 <- lm(lifeExp ~ year + pop, data=kenya)
anova(kenya_model, kenya_model2)
```

```
## Analysis of Variance Table
##
## Model 1: lifeExp ~ year
## Model 2: lifeExp ~ year + pop
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      10 192.035
## 2       9  60.125  1    131.91 19.745 0.001615 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

broom Package: Easier to View Results

- There is a package called broom which makes results of regressions easier to view and compare.
- We will call this package and use the `tidy()` and `glance()` functions.

```
library(broom)
tidy(kenya_model)
glance(kenya_model)
```

broom Package: Easier to View Results

```
##           term      estimate  std.error statistic    p.value
## 1 (Intercept) -356.1009769 145.08537018 -2.454424 0.03399941
## 2          year    0.2065077   0.07329116  2.817634 0.01823454
```

```
##  r.squared adj.r.squared   sigma statistic    p.value df    logLik
## 1 0.4425573    0.386813 4.382174  7.939063 0.01823454  2 -33.66387
##      AIC      BIC deviance df.residual
## 1 73.32775 74.78247 192.0345          10
```

broom Package: Easier to View Results

- We can also compare multiple models at the same time
- Using the commands we learned in data cleaning:

```
tidy1 <- tidy(kenya_model1)
tidy2 <- tidy(kenya_model2)
bind_rows(tidy1, tidy2)
```

broom Package: Easier to View Results

- We can also compare multiple models at the same time
- Using the commands we learned in data cleaning:

##	term	estimate	std.error	statistic	p.value
## 1	(Intercept)	-3.561010e+02	1.450854e+02	-2.454424	0.0339994120
## 2	year	2.065077e-01	7.329116e-02	2.817634	0.0182345440
## 3	(Intercept)	-2.386902e+03	4.649642e+02	-5.133518	0.0006166503
## 4	year	1.250069e+00	2.387936e-01	5.234934	0.0005383790
## 5	pop	-1.918461e-06	4.317401e-07	-4.443555	0.0016151708

broom Package: Easier to View Results

- Again for glance()

```
glance1 <- glance(kenya_model1)
glance2 <- glance(kenya_model2)
bind_rows(glance1, glance2)
```

broom Package: Easier to View Results

```
##   r.squared adj.r.squared   sigma statistic      p.value df   logLik
## 1 0.4425573    0.3868130 4.382174  7.939063 0.0182345440  2 -33.66387
## 2 0.8254669    0.7866818 2.584682 21.283081 0.0003876588  3 -26.69640
##      AIC      BIC  deviance df.residual
## 1 73.32775 74.78247 192.03450         10
## 2 61.39280 63.33243  60.12522          9
```

Variable Selection: Stepwise Regression

```
library(MASS)
fit <- lm(y~x1+x2+x3,data=mydata)
step <- stepAIC(fit, direction="both")
step$anova # display results
```

More Detailed Regression Diagnostics

- We can see more regression diagnostics using the car package
- With this package we have the following functions

More Detailed Regression Diagnostics: Outliers

```
library(car)
outlierTest(kenya_model2) # Bonferonni p-value for most extreme obs
qqPlot(kenya_model2, main="QQ Plot") #qq plot for studentized resid
leveragePlots(kenya_model2) # leverage plots
```


More Detailed Regression Diagnostics: Outliers

```
outlierTest(kenya_model2) # Bonferonni p-value for most extreme obs
```

```
##
## No Studentized residuals with Bonferonni p < 0.05
## Largest |rstudent|:
##      rstudent unadjusted p-value Bonferonni p
## 11 -2.582753      0.032477      0.38972
```

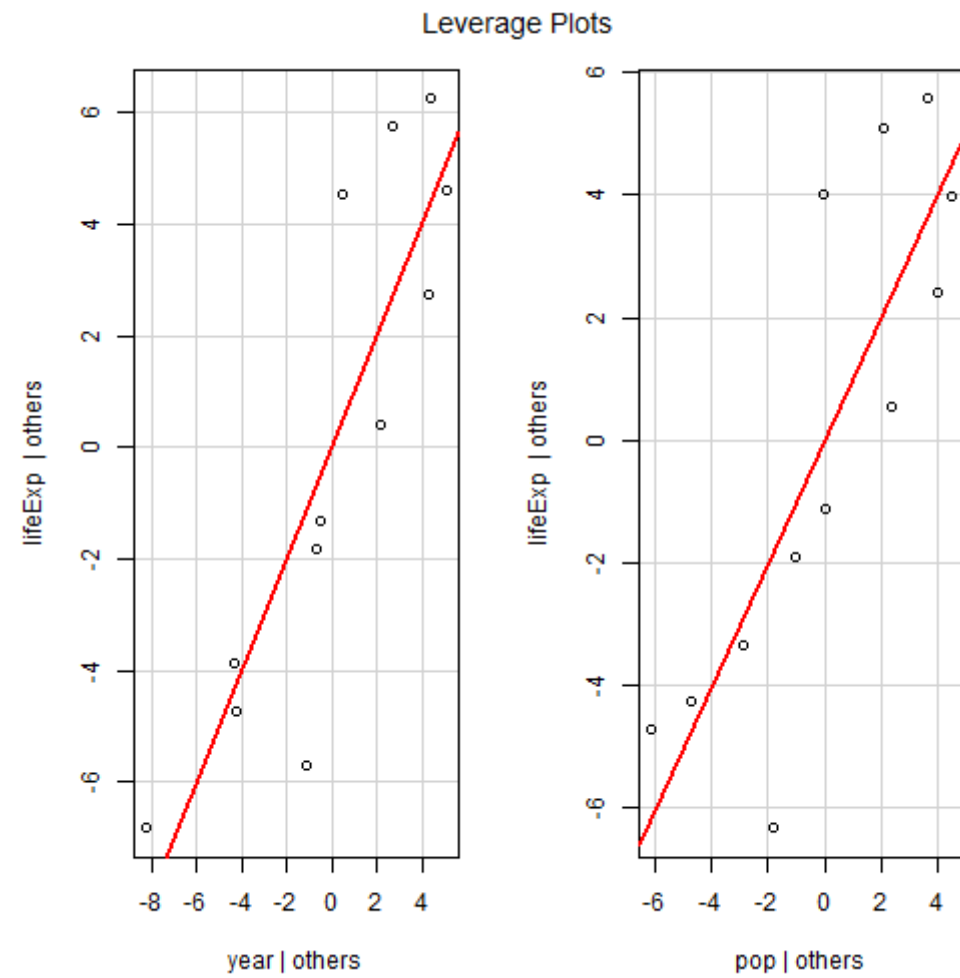
More Detailed Regression Diagnostics: Outliers

```
qqPlot(kenya_model2, main="QQ Plot") #qq plot for studentized resid
```



More Detailed Regression Diagnostics: Outliers

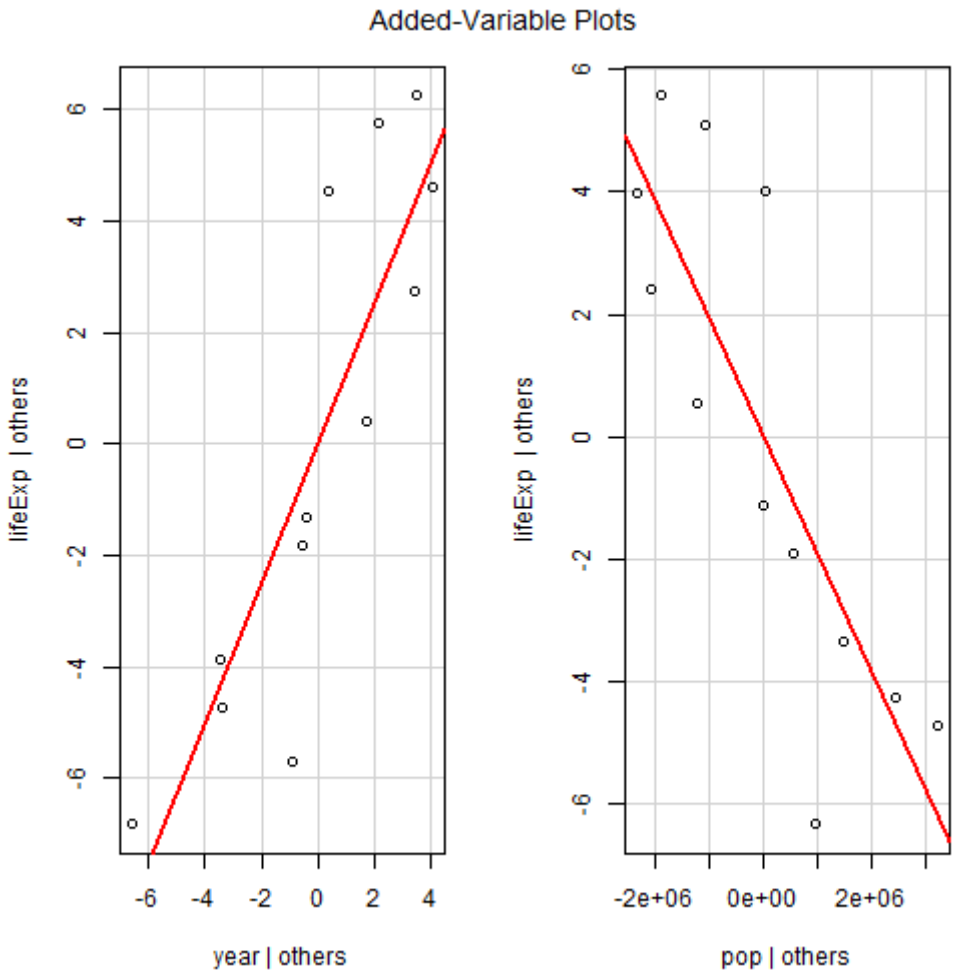
```
leveragePlots(kenya_model2) # leverage plots
```




More Detailed Regression Diagnostics: Influential Observations

```
# Influential Observations
# added variable plots
av.plots(kenya_model2)
# Cook's D plot
# identify D values > 4/(n-k-1)
cutoff <- 4/((nrow(kenya)-length(kenya_model2$coefficients)-2))
plot(kenya_model2, which=4, cook.levels=cutoff)
# Influence Plot
influencePlot(kenya_model2, id.method="identify", main="Influence Plot", sub="Circle size is proportional to Cook's D")
```

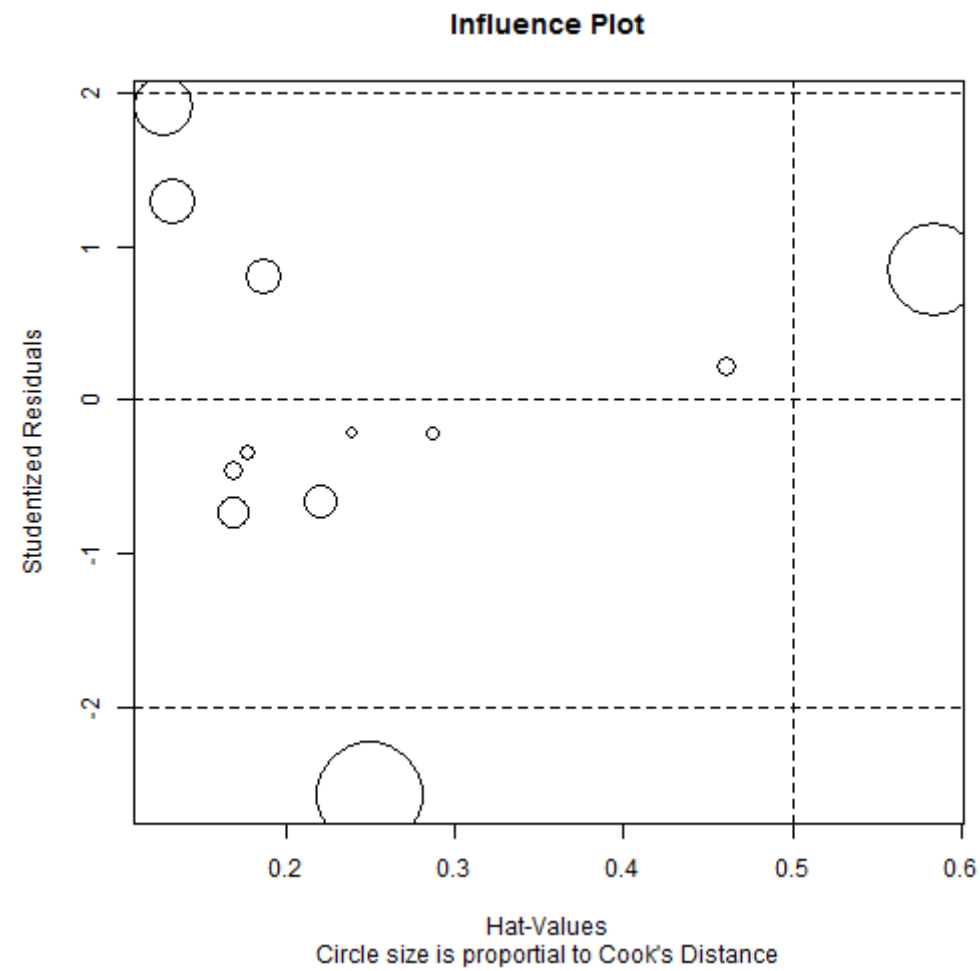
More Detailed Regression Diagnostics: Influential Observations



More Detailed Regression Diagnostics: Influential Observations




More Detailed Regression Diagnostics: Influential Observations




More Detailed Regression Diagnostics: Normality

```
# Normality of Residuals
# qq plot for studentized resid
qqPlot(kenya_model2, main="QQ Plot")
# distribution of studentized residuals
library(MASS)
sresid <- studres(kenya_model2)
hist(sresid, freq=FALSE,
     main="Distribution of Studentized Residuals")
xfit<-seq(min(sresid),max(sresid),length=40)
yfit<-dnorm(xfit)
lines(xfit, yfit)
```

More Detailed Regression Diagnostics: Normality



More Detailed Regression Diagnostics: Normality



More Detailed Regression Diagnostics: Error Variance

```
# Evaluate homoscedasticity
# non-constant error variance test
ncvTest(kenya_model2)
# plot studentized residuals vs. fitted values
spreadLevelPlot(kenya_model2)
```


More Detailed Regression Diagnostics: Error Variance

```
## Non-constant Variance Score Test  
## Variance formula: ~ fitted.values  
## Chisquare = 1.111634    Df = 1    p = 0.2917271
```

More Detailed Regression Diagnostics: Multi-collinearity

<pre># Evaluate Collinearity vif(kenya_model2) # variance inflation factors</pre>
<pre>## year pop ## 30.51456 30.51456</pre>
<pre>sqrt(vif(kenya_model2)) > 2 # problem?</pre>
<pre>## year pop ## TRUE TRUE</pre>

More Detailed Regression Diagnostics: Linearity

```
# Evaluate Nonlinearity  
# component + residual plot  
crPlots(kenya_model2)
```

More Detailed Regression Diagnostics: Linearity

```
## Error in formula(model): object 'fit' not found
```

More Detailed Regression Diagnostics: Autocorrelation

```
# Test for Autocorrelated Errors
durbinWatsonTest(kenya_model2)
```

```
## lag Autocorrelation D-W Statistic p-value
## 1 0.3478423 1.26646 0.024
## Alternative hypothesis: rho != 0
```

More Detailed Regression Diagnostics: Multiple Tests

- We can use one more package to help us with regression diagnostics
- This is the `gvlma` package.

```
library(gvlma)
```

More Detailed Regression Diagnostics: Multiple Tests

```
gvlma(kenya_model)
```

```
##
## Call:
## lm(formula = lifeExp ~ year, data = kenya)
##
## Coefficients:
## (Intercept)      year
##   -356.1010      0.2065
##
##
## ASSESSMENT OF THE LINEAR MODEL ASSUMPTIONS
## USING THE GLOBAL TEST ON 4 DEGREES-OF-FREEDOM:
## Level of Significance =  0.05
##
## Call:
## gvlma(x = kenya_model)
##
##              Value  p-value              Decision
## Global Stat    11.318796 0.023205 Assumptions NOT satisfied!
## Skewness        0.003812 0.950772  Assumptions acceptable.
## Kurtosis        1.047202 0.306152  Assumptions acceptable.
## Link Function    9.764485 0.001779 Assumptions NOT satisfied!
## Heteroscedasticity 0.503297 0.478055  Assumptions acceptable.
```

More Detailed Regression Diagnostics: Multiple Tests

```
gvlma(kenya_model2)
```

```
##
## Call:
## lm(formula = lifeExp ~ year + pop, data = kenya)
##
## Coefficients:
## (Intercept)      year      pop
## -2.387e+03    1.250e+00   -1.918e-06
##
##
## ASSESSMENT OF THE LINEAR MODEL ASSUMPTIONS
## USING THE GLOBAL TEST ON 4 DEGREES-OF-FREEDOM:
## Level of Significance =  0.05
##
## Call:
## gvlma(x = kenya_model2)
##
##
##              Value p-value              Decision
## Global Stat    3.827e+00  0.4300 Assumptions acceptable.
## Skewness       3.116e-05  0.9955 Assumptions acceptable.
## Kurtosis       4.884e-02  0.8251 Assumptions acceptable.
## Link Function  1.643e+00  0.1999 Assumptions acceptable.
## Heteroscedasticity 2.134e+00  0.1440 Assumptions acceptable.
```


Generalized Linear Models in R

Generalized Linear Models in R

- Generalized linear models are fit using the `glm()` function.
- The form of the `glm` function is

```
glm(formula, family=familytype(link=linkfunction), data=)
```

Link Functions

Family	Default Link Function
binomial	(link = "logit")
gaussian	(link = "identity")
Gamma	(link = "inverse")
inverse.gaussian	(link = "1/mu ² ")
poisson	(link = "log")
quasi	(link = "identity", variance = "constant")
quasibinomial	(link = "logit")
quasipoisson	(link = "log")

Further GLM Help

- See `help(glm)` for other modeling options.
- See `help(family)` for other allowable link functions for each family.
- Two subtypes of generalized linear models will be covered here:
 - logistic regression
 - poisson regression

Logistic Regression

- Logistic regression is useful when you are predicting a binary outcome from a set of continuous predictor variables.
- It is frequently preferred over discriminant function analysis because of its less restrictive assumptions.

Fitting Logistic Regression

```
fit <- glm(F~x1+x2+x3,data=mydata,family=binomial
()
)
summary(fit) # display results
confint(fit) # 95% CI for the coefficients
exp(coef(fit)) # exponentiated coefficients
exp(confint(fit)) # 95% CI for exponentiated coefficients
predict(fit, type="response") # predicted values
residuals(fit, type="deviance") # residuals
```

Comparing Nested Logistic Models

- You can use `anova(fit1, fit2, test="Chisq")` to compare nested models.
- Additionally, `cdplot(F~x, data=mydata)` will display the conditional density plot of the binary outcome F on the continuous x variable.

Testing Logistic Regression Models

- We usually determine the goodness of fit for logistic regression based on
 1. **Calibration** - A model is well *calibrated* if the observed and predicted probabilities based on the model are reasonably close.
 2. **Discrimination** - A model has good *discrimination* if the distribution of risk scores for cases and controls separate out. a. This means Cases tend to have higher scores. b. This means Controls tend to have lower scores. c. There is little overlap.

Calibration: Hosmereg-Lemesheu Test

```
library(ResourceSelection)
hoslem.test(sah$dead, fitted(mod.back.auto), g=10)
```

Discrimination: C-Statistic

- We then assess discrimination.
- To do this we use something called **Concordance** or **C Statistic**
- To understand what this is consider 2 different subjects
 1. Subject 1 is dead
 2. Subject 2 is not dead.
- If we consider our model from above it predicts:
 1. \hat{p}_1 the probability that subject 1 is dead.
 2. \hat{p}_2 the probability that subject 2 is dead.

Discrimination: C-Statistic

- The **C Statistic** is given by

$$\Pr(\hat{p}_1 > \hat{p}_2)$$

- If the risk prediction is worthless we find that $C = 0.5$ or essentially the same as flipping a coin.
- If the risk is larger for all who are dead than all who are not dead then we have $C = 1$.
- We typically find this value with a Receiver Operating Characteristic (ROC) curve.

Discrimination: ROC Curve

- Pre-Work for Graph

```
library(ggplot2)
library(ROCR)

prob <- predict(model)
pred <- prediction(prob, data$F)
perf <- performance(pred, "tpr", "fpr")
# I know, the following code is bizarre. Just go with it.
auc <- performance(pred, measure = "auc")
auc <- auc@y.values[[1]]
```

Discrimination: ROC Curve

- Graph

```
roc.data <- data.frame(fpr=unlist(perf@x.values),  
                      tpr=unlist(perf@y.values),  
                      model="GLM")  
ggplot(roc.data, aes(x=fpr, ymin=0, ymax=tpr)) +  
  geom_ribbon(alpha=0.2) + geom_abline(intercept = 0, slope = 1, colour = "gray")+  
  geom_line(aes(y=tpr)) +  
  ggtitle(paste0("ROC Curve w/ AUC=", auc))
```

Poisson Regression

- Poisson regression is useful when predicting an outcome variable representing counts from a set of continuous predictor variables.

Poisson Regression in R

```
# where count is a count and  
# x1-x3 are continuous predictors  
fit <- glm(count ~ x1+x2+x3, data=mydata, family=poisson())  
summary(fit) display results
```

Poisson Regression with Overdispersion

- If you have overdispersion you may want to use `quasipoisson()` instead of `poisson()`.